



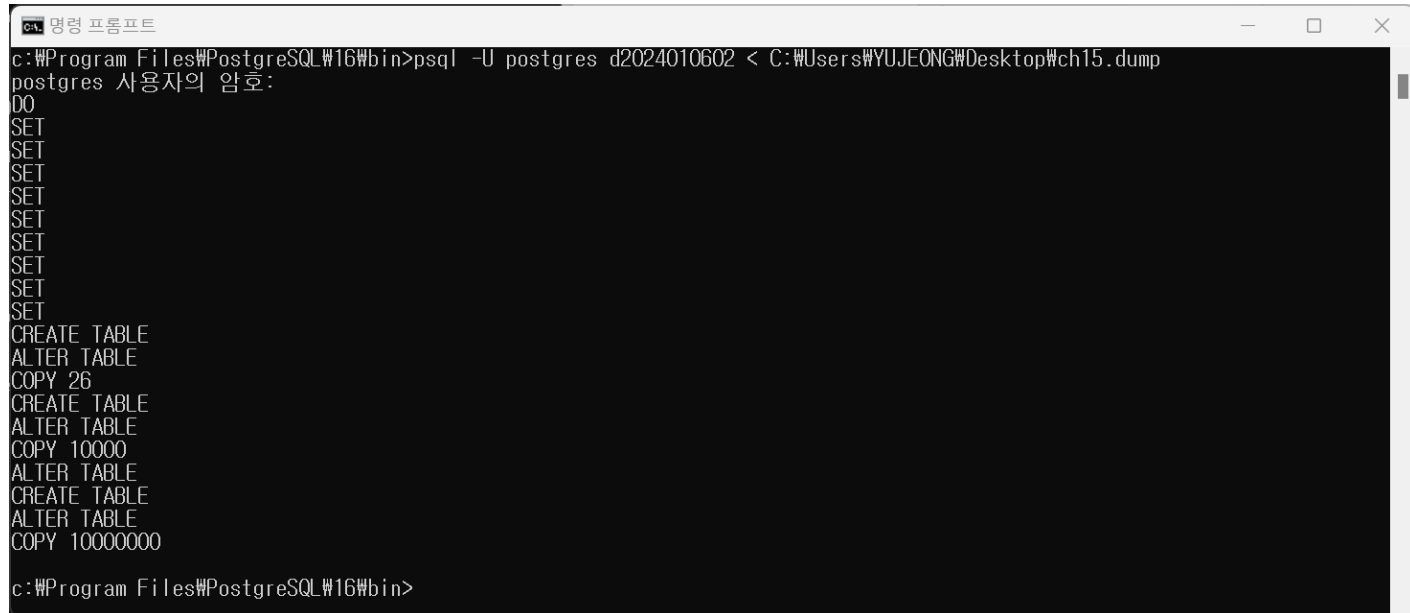
KOREA UNIVERSITY
DATABASE LAB

Chapter 15 - Lab

Query Processing

Lab Setup (Windows)

- Open **Command Prompt (cmd.exe)** and type the following commands:
 1. `cd C:\Program Files\PostgreSQL\16\bin`
 - This is the **default** path. If you installed it somewhere else, go to that path.
 2. `psql -U postgres d{StudentID} < [filepath]\ch15.dump`
 - For **[filepath]**, type the path where you downloaded “ch15.dump”.
 3. Type your own PostgreSQL password



```
cmd - 명령 프롬프트
c:\Program Files\PostgreSQL\16\bin>psql -U postgres d2024010602 < C:\Users\WYUJEONG\Desktop\ch15.dump
postgres 사용자의 암호:
DO
SET
SET
SET
SET
SET
SET
SET
SET
SET
SET
CREATE TABLE
ALTER TABLE
COPY 26
CREATE TABLE
ALTER TABLE
COPY 10000
ALTER TABLE
CREATE TABLE
ALTER TABLE
COPY 10000000
c:\Program Files\PostgreSQL\16\bin>
```

Lab Setup (Mac OS X)

- Open **Terminal** and type the following commands:
 1. `cd /Library/PostgreSQL/16/bin`
 - This is the **default** path. If you installed it somewhere else, go to that path.
 2. `./psql -U postgres d{StudentID} < [filepath]/ch15.dump`
 - For **[filepath]**, type the path where you downloaded “ch15.dump”.
 3. Type your own PostgreSQL password

Lab Setup

- Execute PostgreSQL **SQL Shell (psql)** and login your database
 - Server [localhost]: Press the enter key
 - Database [postgres]: Press the enter key
 - Port [5432]: Press the enter key
 - Username [postgres]: Press the enter key
 - Password for user postgres: **Type your own password**
 - **\c d{StudentID}**
- Type on psql command line
 - SET enable_bitmapscan=false;
 - SET max_parallel_workers_per_gather=0;

Table Information

- Schema of each table is as follows:

“table1”
(10,000,000 rows)

| Attribute | Data Type | Cardinality | Features |
|-----------|---------------|-------------|-------------|
| sorted | integer | 2,000,000 | Sorted |
| unsorted | integer | 1,986,519 | Unsorted |
| rndm | integer | 100,000 | Dummy field |
| dummy | character(40) | 1 | Dummy field |

“nation”
(26 rows)

| Attribute | Data Type | Cardinality | Features |
|--------------|-----------------------|-------------|----------|
| n_nationkey | integer | 26 | |
| n_nationname | character varying(50) | 26 | |

Table Information

- Schema of each table is as follows:

“supplier” (10,000 rows)

| Attribute | Data Type | Cardinality | Features |
|-------------|------------------------|-------------|-------------|
| s_suppkey | integer | 10,000 | primary key |
| s_name | character varying(200) | 10,000 | |
| s_address | character varying(200) | 10,000 | |
| s_nationkey | integer | 25 | |
| s_phone | character varying(200) | 10,000 | |
| s_acctbal | double precision | 9,955 | |
| s_comment | character varying(200) | 10,000 | |

Exercise 1

- Consider two join cases below:
 - a. Equi-join two tables of “supplier” and “nation” with “s_nationkey” and “n_nationkey” as join keys
 - b. Attach the above SQL statement with “ORDER BY s_nationkey”, and test again
 - Hint: ORDER BY s_nationkey
- Estimate **join algorithms** that are applied to each case
- Verify whether your estimation is correct by using ‘EXPLAIN ANALYZE’

Exercise 2

- Consider two self-join cases below:
 - a. Equi-join of the table “nation” with “n_nationkey” as a join key
 - b. Non equi-join of the table “nation” with “n_nationkey” as a join key
- Estimate **join algorithms** that are applied to each case
- Verify whether your estimation is correct by using ‘EXPLAIN ANALYZE’

Exercise 3

- Consider two join cases below:
 - a. Equi-join two tables of “supplier” and “table1” with “s_suppkey” and “sorted” as join keys
 - b. Equi-join two tables of “supplier” and “table1” with “s_suppkey” and “unsorted” as join keys
- Estimate **execution time** and **join algorithms** that are applied to each case
- Verify whether your estimation is correct by using ‘EXPLAIN ANALYZE’

Exercise 4

- Create three indexes:
 - CREATE INDEX sorted_idx on table1(sorted);
 - CREATE INDEX unsorted_idx on table1(unsorted);
 - CREATE INDEX suppkey_idx on supplier(s_suppkey);
 - ANALYZE table1;
 - ANALYZE supplier;
- Consider two join cases below:
 - a. Equi-join two tables of “supplier” and “table1” with “s_suppkey” and “sorted” as join keys
 - b. Equi-join two tables of “supplier” and “table1” with “s_suppkey” and “unsorted” as join keys
- Estimate **execution time** and **join algorithms** that are applied to each case
- Verify whether your estimation is correct by using ‘EXPLAIN ANALYZE’

Exercise 5-1

- Type on psql command line
 - SET enable_memoize = off;
- Consider two join cases below:
 - a. Equi-join two tables of “supplier” and “table1” with “s_nationkey” and “sorted” as join keys
 - b. Equi-join two tables of “supplier” and “table1” with “s_nationkey” and “unsorted” as join keys
- Estimate **execution time** and **join algorithms** that are applied to each case
- Verify whether your estimation is correct by using ‘EXPLAIN ANALYZE’

Exercise 5-2

- Add a large number of records to the s_nationkey as follows.
 - INSERT INTO supplier(s_suppkey, s_nationkey) SELECT generate_series(10001, 10000000), generate_series(10001, 10000000);
 - ANALYZE supplier;
- Consider two join cases below:
 - a. Equi-join two tables of “supplier” and “table1” with “s_nationkey” and “sorted” as join keys
 - b. Equi-join two tables of “supplier” and “table1” with “s_nationkey” and “unsorted” as join keys
- Estimate **execution time** and **join algorithms** that are applied to each case
- Verify whether your estimation is correct by using ‘EXPLAIN ANALYZE’

Homework

- Complete today's practice exercises
- Write your queries and take screenshots of execution results
- Submit your report on blackboard
 - 10:29:59, October 21th, 2024
 - **Only PDF files** are accepted
 - **No late submission**

End of Lab