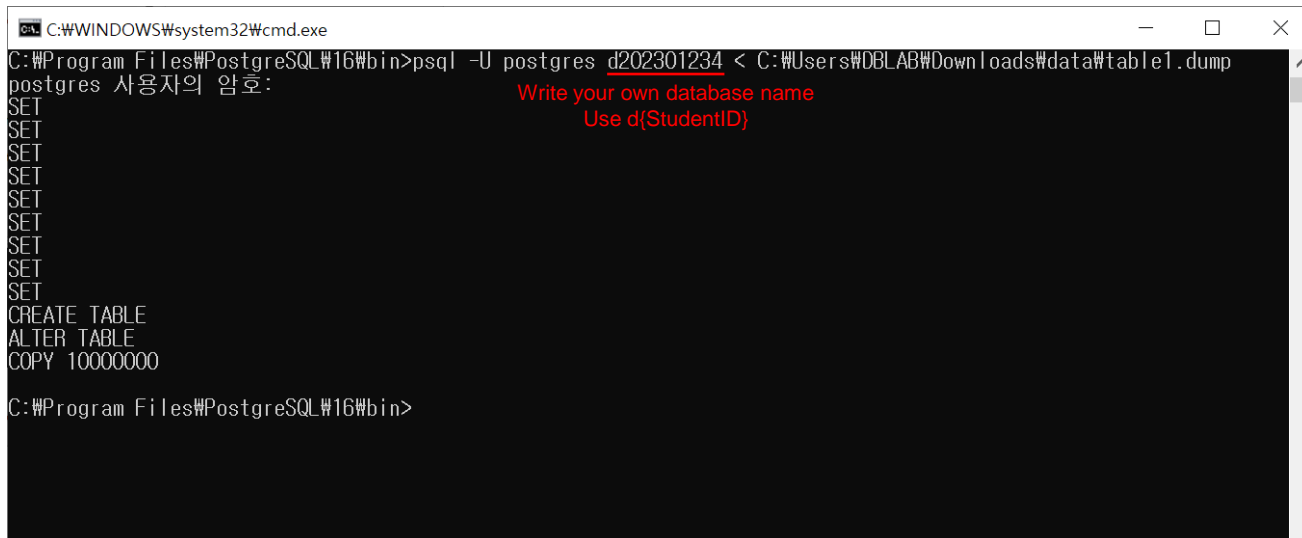


Chapter 14 - Lab

Indexing 1

Lab Setup (Windows)

- Download the “table1.dump” file from blackboard
- Open **Command Prompt (cmd.exe)** and type the following commands:
 1. `cd C:\Program Files\PostgreSQL\16\bin`
 - This is the **default** path. If you installed it somewhere else, go to that path.
 2. `psql -U postgres d{StudentID} < [filepath]\table1.dump`
 - For [filepath], type the path where you downloaded “table1.dump”.
 3. Type your own PostgreSQL password



```
C:\WINDOWS\system32\cmd.exe
C:\Program Files\PostgreSQL\16\bin>psql -U postgres d202301234 < C:\Users\DBLAB\Downloads\data\table1.dump
postgres 사용자의 암호:
SET
SET
SET
SET
SET
SET
SET
SET
SET
SET
SET
CREATE TABLE
ALTER TABLE
COPY 10000000
C:\Program Files\PostgreSQL\16\bin>
```

Write your own database name
Use d{StudentID}

Lab Setup (Max OS X)

- Download the “table1.dump” file from blackboard
- Open **Terminal** and type the following commands:
 1. `cd /Library/PostgreSQL/16/bin`
 - This is the **default** path. If you installed it somewhere else, go to that path.
 2. `psql -U postgres d{StudentID} < [filepath]\table1.dump`
 - For [filepath], type the path where you downloaded “table1.dump”.
 3. Type your own PostgreSQL password



```
bin — zsh — 80x24
Last login: Thu Sep 22 16:32:04 on ttys000
[(base) hyubjinlee@hyubjinleeui-MacBookPro ~ % cd /Library/PostgreSQL/14/bin ]
[(base) hyubjinlee@hyubjinleeui-MacBookPro bin % ./psql -U postgres postgres < /Users/hyubjinlee/Desktop/table1.dump
Password for user postgres:
SET
SET
SET
SET
SET
SET
SET
SET
SET
SET
CREATE TABLE
ALTER TABLE
COPY 10000000
(base) hyubjinlee@hyubjinleeui-MacBookPro bin %
```

Write your own database name
Use d{StudentID}

Lab Setup

- Execute PostgreSQL **SQL Shell (psql)** and login your database
 - Server [localhost]: Press the enter key
 - Database [postgres]: Press the enter key
 - Port [5432]: Press the enter key
 - Username [postgres]: Press the enter key
 - Password for user postgres: **Type your own password**
 - **\c d{StudentID}**

```
postgres=# \c d202301234
접속정보: 데이터베이스="d202301234", 사용자="postgres".
d202301234=#
```

Your answers must be displayed along with your student ID.

- Type on psql command line
 - SET enable_bitmapscan=false;
 - SET max_parallel_workers_per_gather=0;

Table Information

- “table1” has 10,000,000 rows
- “table1”'s schema is as follows:

Attribute	Data Type	Cardinality	Features
sorted	integer	2,000,000	Sorted
unsorted	integer	1,986,519	Unsorted
rndm	integer	100,000	Dummy field
dummy	character(40)	1	Dummy field

Exercise 1

- Create two indexes on “table1”
 - Indexed attributes are “sorted” and “unsorted”
 - Type “\h CREATE INDEX” for detailed index creation syntax

Exercise 2

- PostgreSQL supports following index-based query execution plans
 - Seq scan: All rows in a table are read sequentially
 - Index scan: Some (or all) rows in a table are read after traversing an index
 - Index only scan: Query is processed in an index, not accessing table data
- 'EXPLAIN ANALYZE' statement shows the query plan and execution time of the query
 - E.g., `> EXPLAIN ANALYZE SELECT * FROM table1;`

Exercise 2

- a. Make (and execute) three queries each of which uses *seq scan*, *index scan*, and *index only scan* respectively
- b. Make two queries that are expected to use indices on attributes “sorted” and “unsorted” respectively, then compare their execution times
- c. Execute your queries to Exercise 2.b. after executing each of the following queries respectively, then compare their execution times
 - `CLUSTER table1 USING idx_unsorted;`
 - `CLUSTER table1 USING idx_sorted;`
- d. Execute and compare the following two queries:
 - `SELECT sorted, rndm FROM table1 WHERE sorted>1999231 AND rndm=1005;`
 - `SELECT sorted, rndm FROM table1 WHERE sorted<1999231 AND rndm=1005;`
 - Explain why their query plans are different

Exercise 3

- Setup: Create a synthetic data set that has 5,000,000 rows
 - CREATE TABLE pool (val integer);
 - INSERT INTO pool(val) SELECT * FROM (SELECT generate_series(1,5000000)) as T;
 - Type on psql command line
 - SET enable_bitmapscan=false; \timing
- Consider two cases below. Which case will take a longer time?
 1. Inserting tuples in a table, and then creating index
 2. Creating index, and then inserting tuples in a table
- Compare the execution time t_1 and t_2
 - $t_1 = t_{1.insert} + t_{1.create_index}$
 - Tuple insertion → Index creation
 - $t_2 = t_{2.create_index} + t_{2.insert}$
 - Index creation → Tuple insertion

Exercise 3 - Hints

- Create an empty table named “table10” and “table20”
 - CREATE TABLE table10 (val integer);
 - CREATE TABLE table20 (val integer);
- Use “table10” to measure t_1 and “table20” to measure t_2
 - Inserting tuples into tables
 - Utilize the “pool” table for a synthetic dataset
 - Insert all the tuples of “pool” into “table10” and “table20”
 - E.g., > INSERT INTO table10 (SELECT * FROM pool);

Homework

- Complete today's practice exercises
- Write your queries and take screenshots of execution results
- Submit your report on blackboard
 - 10:29, October 1st, 2024
 - **Only PDF files** are accepted
 - **No late submission**

End of Lab