

Chapter 14: Indexing (2/2)

(+ Ch. 24)

Outline

- Basic Concepts
- Ordered Indices
- B⁺-Tree Index
- Hashing
- Write-optimized indices
- Spatio-Temporal Indexing



Chapter 24.

Hashing

Static Hashing

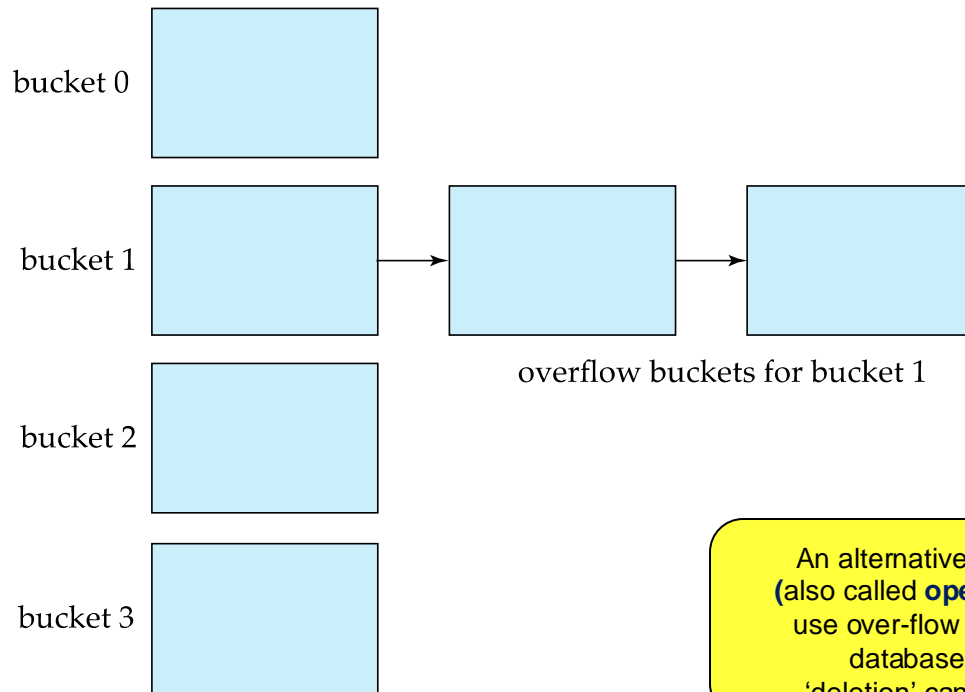
- A **bucket** is a unit of storage containing one or more entries (a bucket is typically a disk block).
 - We obtain the bucket of an entry from its search-key value using a **hash function**.
- Hash function h is a function from the set of all search-key values K to the set of all bucket addresses B .
- Hash function is used to locate entries for access, insertion as well as deletion.
- Entries with different search-key values may be mapped to the same bucket; thus entire bucket has to be searched sequentially to locate an entry.
- In a **hash index**, buckets store entries with pointers to records
- In a **hash file-organization** buckets store records

Handling of Bucket Overflows

- Bucket overflow can occur because of
 - Insufficient buckets
 - Skew in distribution of records because of the following two reasons:
 1. chosen hash function produces non-uniform distribution of key values
 2. multiple records have same search-key value
- Although the probability of bucket overflow can be reduced, it cannot be eliminated; it is handled by using **overflow buckets**.

Handling of Bucket Overflows (Cont.)

- **Overflow chaining** – the overflow buckets of a given bucket are chained together in a linked list.
- Above scheme is called **closed addressing** (also called **closed hashing**)



An alternative, called **open addressing** (also called **open hashing**) which does not use over-flow buckets, is not suitable for database applications because 'deletion' cannot be handled efficiently.

Example of Hash File Organization

Hash file organization of *instructor* file, using *dept_name* as key

- There are 8 buckets,
- The i^{th} character is computed as integer i .
 - $a \rightarrow 1, \dots, z \rightarrow 26$
- The hash function returns the sum of the converted values of the characters modulo 8
 - E.g. $h(\text{Music}) = 1$ $h(\text{History}) = 2$
 $h(\text{Physics}) = 3$ $h(\text{Elec. Eng.}) = 3$

Example of Hash File Organization

Hash file organization of *instructor* file, using *dept_name* as key.

bucket 0

bucket 1

15151	Mozart	Music	40000

bucket 2

32343	El Said	History	80000
58583	Califieri	History	60000

bucket 3

22222	Einstein	Physics	95000
33456	Gold	Physics	87000
98345	Kim	Elec. Eng.	80000

bucket 4

12121	Wu	Finance	90000
76543	Singh	Finance	80000

bucket 5

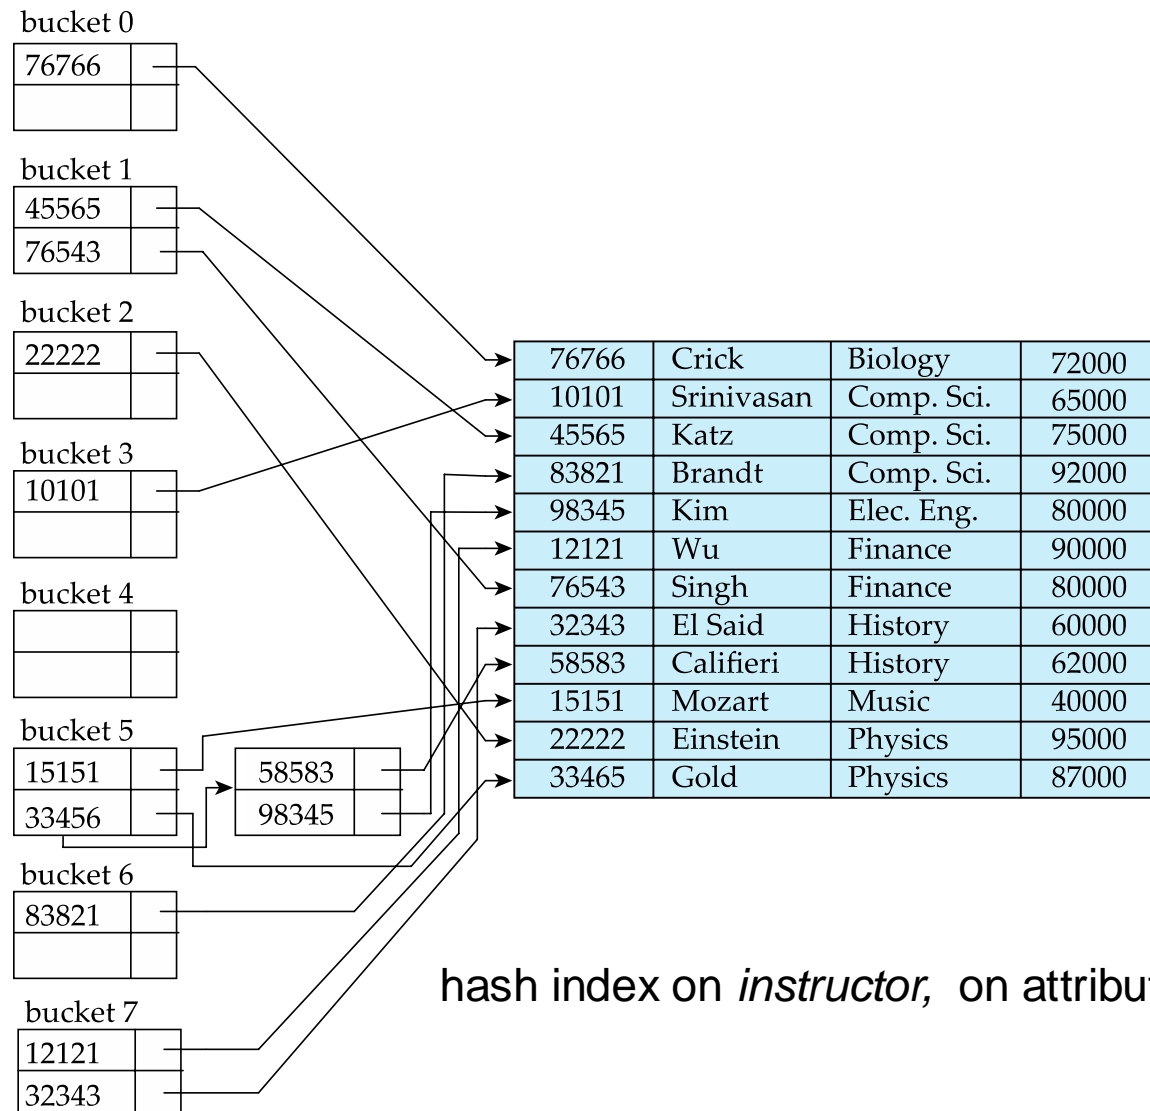
76766	Crick	Biology	72000

bucket 6

10101	Srinivasan	Comp. Sci.	65000
45565	Katz	Comp. Sci.	75000
83821	Brandt	Comp. Sci.	92000

bucket 7

Example of Hash Index



hash index on *instructor*, on attribute *ID*

Deficiencies of Static Hashing

- In static hashing, function h maps search-key values to a fixed set of B of bucket addresses. Databases grow or shrink with time.
 - If initial number of buckets is too small, and file grows, performance will degrade due to too much overflows.
 - If space is allocated for anticipated growth, a significant amount of space will be wasted initially (and buckets will be underfull).
 - If database shrinks, again space will be wasted.
- One solution: periodic re-organization of the file with a new hash function
 - If number of entries in a hash table becomes (say) 1.5 times size of hash table,
 - create new hash table of size (say) 2 times the size of the previous hash table
 - Rehash all entries to new table
 - Expensive, disrupts normal operations
- Better solution: allow the number of buckets to be modified **dynamically**.

Dynamic Hashing

- (See Ch. 24)
- Linear Hashing
 - Do rehashing in an incremental manner
- Extendable Hashing
 - Tailored to disk-based hashing, with buckets shared by multiple hash values
 - Doubling of # of entries in hash table, without doubling # of buckets

Comparison of Ordered Indexing and Hashing

- Cost of periodic re-organization
- Relative frequency of insertions and deletions
- Is it desirable to optimize average access time at the expense of worst-case access time?
- Expected type of queries:
 - Hashing is generally better at retrieving records having a specified value of the key.
 - If range queries are common, ordered indices are to be preferred
- In practice:
 - PostgreSQL supports hash indices, but discourages use due to poor performance
 - Oracle supports static hash organization, but not hash indices
 - SQLServer supports only B⁺-trees

Multiple (Single-Attribute) Indices

- Use multiple indices for certain types of queries.
- Example:
select *ID*
from *instructor*
where *dept_name* = “Finance” **and** *salary* = 80000
- Possible strategies for processing query using indices on single attributes:
 1. Use index on *dept_name* to find instructors with department name Finance; test *salary* = 80000
 2. Use index on *salary* to find instructors with a salary of \$80000; test *dept_name* = “Finance”.
 3. Use *dept_name* index to find pointers to all records pertaining to the “Finance” department. Similarly use index on *salary*. Take intersection of both sets of pointers obtained.

Indices on Multiple Search Keys

- **Composite search keys** are search keys containing more than one attribute
 - E.g., (*dept_name*, *salary*)
- Lexicographic ordering: $(a_1, a_2) < (b_1, b_2)$ if either
 - $a_1 < b_1$, or
 - $a_1 = b_1$ and $a_2 < b_2$

Indices on Multiple Search Keys (cont'd)

Suppose we have an index on combined search-key
(*dept_name*, *salary*).

- With the **where** clause
 where *dept_name* = “Finance” **and** *salary* = 80000
the index on (*dept_name*, *salary*) can be used to fetch only records that satisfy both conditions.
 - Using separate indices is less efficient — we may fetch many records (or pointers) that satisfy only one of the conditions.
- Can also efficiently handle
 where *dept_name* = “Finance” **and** *salary* < 80000
- But cannot efficiently handle
 where *dept_name* < “Finance” **and** *salary* = 80000
 - May fetch many records that satisfy the first but not the second condition

Other Features

■ Covering indices

- Add extra attributes to index so (some) queries can be answered without fetching the actual records
(e.g.) index on 'name' (+ 'salary' values on the leaf index entries)
- Store extra attributes only at leaf
 - Why?
- Particularly useful for secondary indices
 - Why?

Index Definition in SQL

- Create an index

create index <index-name> **on** <relation-name>
(<attribute-list>)

E.g.,: **create index** *b-index* **on** *branch(branch_name)*

- Use **create unique index** to indirectly specify and enforce the condition that the search key is a candidate key is a candidate key.
 - Not really required if SQL **unique** integrity constraint is supported

- Drop an index

drop index <index-name>

Creation of Indices

- Example

```
create index instructor_dept_index on instructor (dept_name) ;  
drop index instructor_dept_index ;
```

- Most database systems allow specification of type of index (and clustering).
- Indices on primary key created automatically by most databases
 - Why?
- Some database also create indices on foreign key attributes automatically
 - Why might such an index be useful for this query ?
 - $takes \bowtie \sigma_{name='Shankar'} (student)$
- Indices can greatly speed up lookups, but impose cost on updates
 - Index tuning assistants/wizards supported on several databases to help choose indices, based on query and update workload

Write Optimized Indices

- Performance of B⁺-trees can be poor for write-intensive workloads
 - One I/O per leaf, although assuming all internal nodes are in memory
 - With magnetic disks, < 100 inserts per second per disk
 - With flash memory, one page overwrite per insert (+ page erase ?)
- Two approaches to reducing cost of writes
 - **Log-structured merge tree**
 - **Buffer tree**

Bitmap Indices

- Bitmap indices are a special type of index designed for efficient querying on multiple keys
- Records in a relation are assumed to be numbered sequentially from, say, 0
 - Given a number n , it must be easy to retrieve record n
 - Particularly easy if records are of fixed size
- Applicable on attributes that take on a relatively small number of distinct values
 - E.g., gender, country, state, ...
 - E.g., income-level (income broken up into a small number of levels such as 0-9999, 10000-19999, 20000-50000, 50000- infinity)
- A bitmap is simply an array of bits

Bitmap Indices (Cont.)

- In its simplest form, a bitmap index on an attribute has a bitmap for each value (possibly including the *null*) of the attribute
 - Bitmap has as many bits as records
 - In a bitmap for value *v*, the bit for a record is 1 if the record has the value *v* for the attribute, and is 0 otherwise
- Example

record number	ID	gender	income_level
0	76766	m	L1
1	22222	f	L2
2	12121	f	L1
3	15151	m	L4
4	58583	f	L3

Bitmaps for *gender*

m	10010
f	01101

Bitmaps for *income_level*

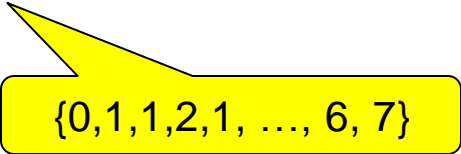
L1	10100
L2	01000
L3	00001
L4	00010
L5	00000

Bitmap Indices (Cont.)

- Bitmap indices are useful for queries on multiple attributes
 - not particularly useful for single attribute queries
- Queries are answered using bitmap operations
 - Intersection (and)
 - Union (or)
- Each operation takes two bitmaps of the same size and applies the operation on corresponding bits to get the result bitmap
 - E.g., $100110 \text{ AND } 110011 = 100010$
 $100110 \text{ OR } 110011 = 110111$
 $\text{NOT } 100110 = 011001$
 - Males with income level L1: $10010 \text{ AND } 10100 = 10000$
 - Can then retrieve required tuples.
 - Counting number of matching tuples is even faster
- Bitmap indices generally very small compared with relation size

Efficient Implementation of Bitmap Operations

- Bitmaps are packed into words; a bit-wise **and** (a basic CPU instruction) computes **and** of 32 or 64 bits at once
 - E.g., 1-million-bit maps can be **and**-ed with just 31,250 instruction
- Counting number of 1s can be done fast by a trick:
 - Use each byte to index into a precomputed array of 256 ($=2^8$) elements each storing the count of 1s in the binary representation
 - Can use pairs of bytes to speed up further at a higher memory cost
 - Add up the retrieved counts
- Bitmaps can be used instead of Tuple-ID lists at leaf levels of B⁺-trees, for values that have a large number of matching records
 - Worthwhile if $> 1/64$ of the records have that value, assuming a tuple-id is 64 bits
 - $64 \text{ bits} * N \text{ records} * \text{selectivity}$ vs. $N \text{ bits}$
 - Above technique merges benefits of bitmap and B⁺-tree indices



{0,1,1,2,1, ..., 6, 7}


Spatial and Temporal Indices

Spatial Data

- Databases can store data types such as lines, polygons, in addition to raster images
 - Allows relational databases to store and retrieve spatial information.
 - Queries can use spatial conditions (e.g. contains or overlaps).
 - Queries can mix spatial and nonspatial conditions.
- **Nearest neighbor queries**, given a point or an object, find the nearest object that satisfies given conditions.
- **Range queries** deal with spatial regions. e.g., ask for objects that lie partially or fully inside a specified region.
- Queries that compute intersections or **unions** of regions.
- **Spatial join** of two spatial relations with the location playing the role of join attribute.

K-d trees


- **k-d tree** - early structure used for indexing spatial points in multiple dimensions.
- Each level of a *k-d* tree partitions the space into two.
 - Choose one dimension for partitioning at the root level of the tree.
 - Choose another dimensions for partitioning in nodes at the next level and so on, cycling through the dimensions.
- In each node, approximately half of the points stored in the sub-tree fall on one side and half on the other.
- Partitioning stops when a node has less than a given number of points.



- The **k-d-B tree** extends the *k-d* tree to allow multiple child nodes for each internal node; well-suited for secondary storage.

Quadrees

- Each node of a quadtree is associated with a rectangular region of space; the top node is associated with the entire target space.
- Each non-leaf nodes divides its region into four equal sized quadrants
 - correspondingly each such node has four child nodes corresponding to the four quadrants and so on
- Leaf nodes have between zero and some fixed maximum number of points (set to 1 in the below example).

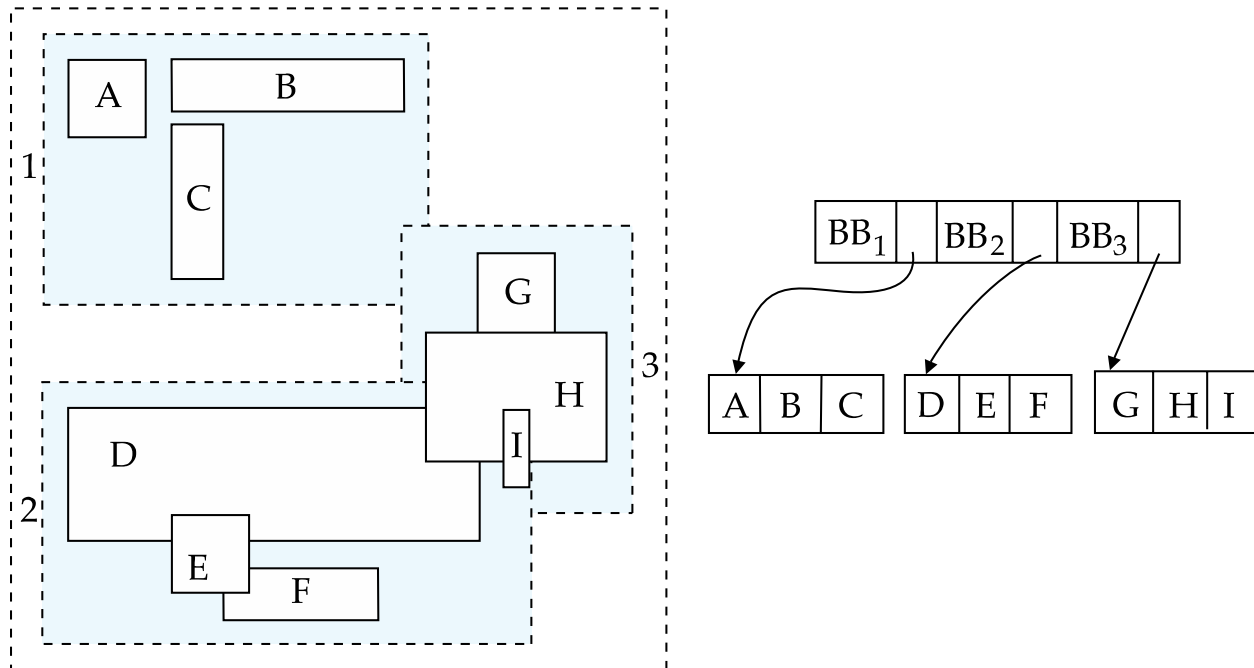


R-Trees

- **R-trees** are a N-dimensional extension of B⁺-trees, useful for indexing sets of rectangles and other polygons.
- Supported in many modern database systems, along with variants like R⁺ - trees and R^{*}-trees.
- Basic idea:
 - Generalize the notion of a one-dimensional interval associated with each B⁺-tree node to an N-dimensional interval, that is, an N-dimensional rectangle.
- Will consider only the two-dimensional case ($N = 2$)
 - generalization for $N > 2$ is straightforward, although R-trees work well only for relatively small N
- The **bounding box** of a node is a minimum sized rectangle that contains all the rectangles/polygons associated with the node
 - *Bounding boxes of children of a node are allowed to overlap*

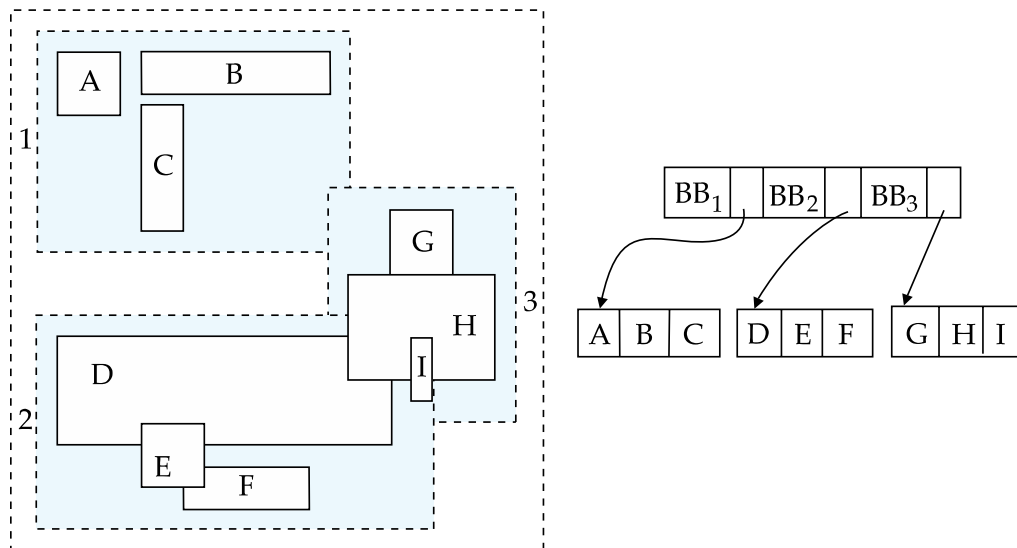
Example R-Tree

- A set of rectangles (solid line) and the bounding boxes (dashed line) of the nodes of an R-tree for the rectangles.
- The R-tree is shown on the right.



Search in R-Trees

- To find data items intersecting a given query point/region, do the following, starting from the root node:
 - If the node is a leaf node, output the data items whose keys intersect the given query point/region.
 - Else, for each child of the current node whose bounding box intersects the query point/region, recursively search the child
- Can be very inefficient in worst case since multiple paths may need to be searched, but works acceptably in practice.



Indexing Temporal Data

- Temporal data refers to data that has an associated time period (interval)
 - Example: a temporal version of the *course* relation

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>start</i>	<i>end</i>
BIO-101	Intro. to Biology	Biology	4	1985-01-01	9999-12-31
CS-201	Intro. to C	Comp. Sci.	4	1985-01-01	1999-01-01
CS-201	Intro. to Java	Comp. Sci.	4	1999-01-01	2010-01-01
CS-201	Intro. to Python	Comp. Sci.	4	2010-01-01	9999-12-31

- Time interval has a start and end time
 - End time set to infinity (or large date such as 9999-12-31) if a tuple is currently valid and its validity end time is not currently known
- Query may ask for all tuples that are valid at a point in time or during a time interval
 - Index on valid time period speeds up this task

Indexing Temporal Data (Cont.)

- How to retrieve records with attribute a is v and valid at time t
 1. Use index created on a , then filter records using time values
 2. Use spatial index, such as R-tree, with attribute a as one dimension, and time as another dimension
 - Valid time forms an *interval* in the time dimension
 - Tuples that are currently valid cause problems, since value is infinite or very large (→ not good for spatial indices)
 - Solution → Store all current tuples (with end time as infinity) in a separate index, indexed on $(a, \text{start-time})$ using B⁺-tree
 - To find tuples valid at a point in time t , search for tuples in the range $(a = v, \text{start-time} \leq t)$ in the *current tuple index* (e.g. B⁺-tree) + search for tuples in the range $(a = v, 0 \text{ start-time} \leq t \leq \text{end-time})$ in the spatial index (e.g. R-tree)
- Temporal index on primary key can help enforce temporal primary key constraint (i.e., non-overlapping time intervals with the same PK).

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>start</i>	<i>end</i>
BIO-101	Intro. to Biology	Biology	4	1985-01-01	9999-12-31
CS-201	Intro. to C	Comp. Sci.	4	1985-01-01	1999-01-01
CS-201	Intro. to Java	Comp. Sci.	4	1999-01-01	2010-01-01
CS-201	Intro. to Python	Comp. Sci.	4	2010-01-01	9999-12-31

End of Chapter 14 (+Ch. 24)