

Lecture 24: Visualization Pipeline

Dec 5, 2024

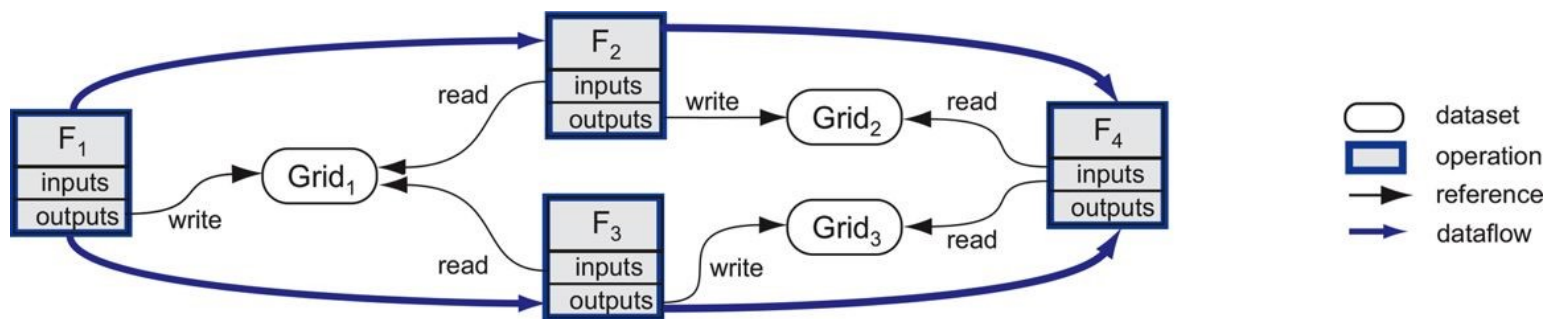
Won-Ki Jeong

(wkjeong@korea.ac.kr)



Visualization Pipeline

- **A dataflow network** in which computation is described as a collection of executable **modules** that are connected in a **directed graph** representing data movement between modules



Modules

- Basic building blocks (functional units) of a pipeline
- Encapsulating algorithms
- Having generic connection ports
 - Inputs and outputs
- Interchangeable
 - As long as input and output data is compatible
- Node in a directed graph




Types of Modules

- Source
 - Producing data
 - File reader, synthetic data generator
- Sink
 - Accept data as an input and no further result
 - File writers, rendering modules
- Filter
 - At least one input and one output
 - Transform data



The Simplest Pipeline Example




Connections

- Directional attachments from the output port of one module to the input port of another module
- Arc in a directed graph



Visualization Pipeline with Branches



Execution Management

- Determine how and when modules get executed
 - Who starts (drives) execution?
 - Do we need to store intermediate results?
 - Where to control execution?
 - What if the data is too big to fit to a module?




Execution Drivers: Event-driven

- Launch execution as data become available in source module
- Source module pushes the data to downstream modules and triggers an event to execute them
 - Execution is initiated from source
 - Called *push model*
- ex) time-varying data visualization



Example: Time-varying Volume Rendering

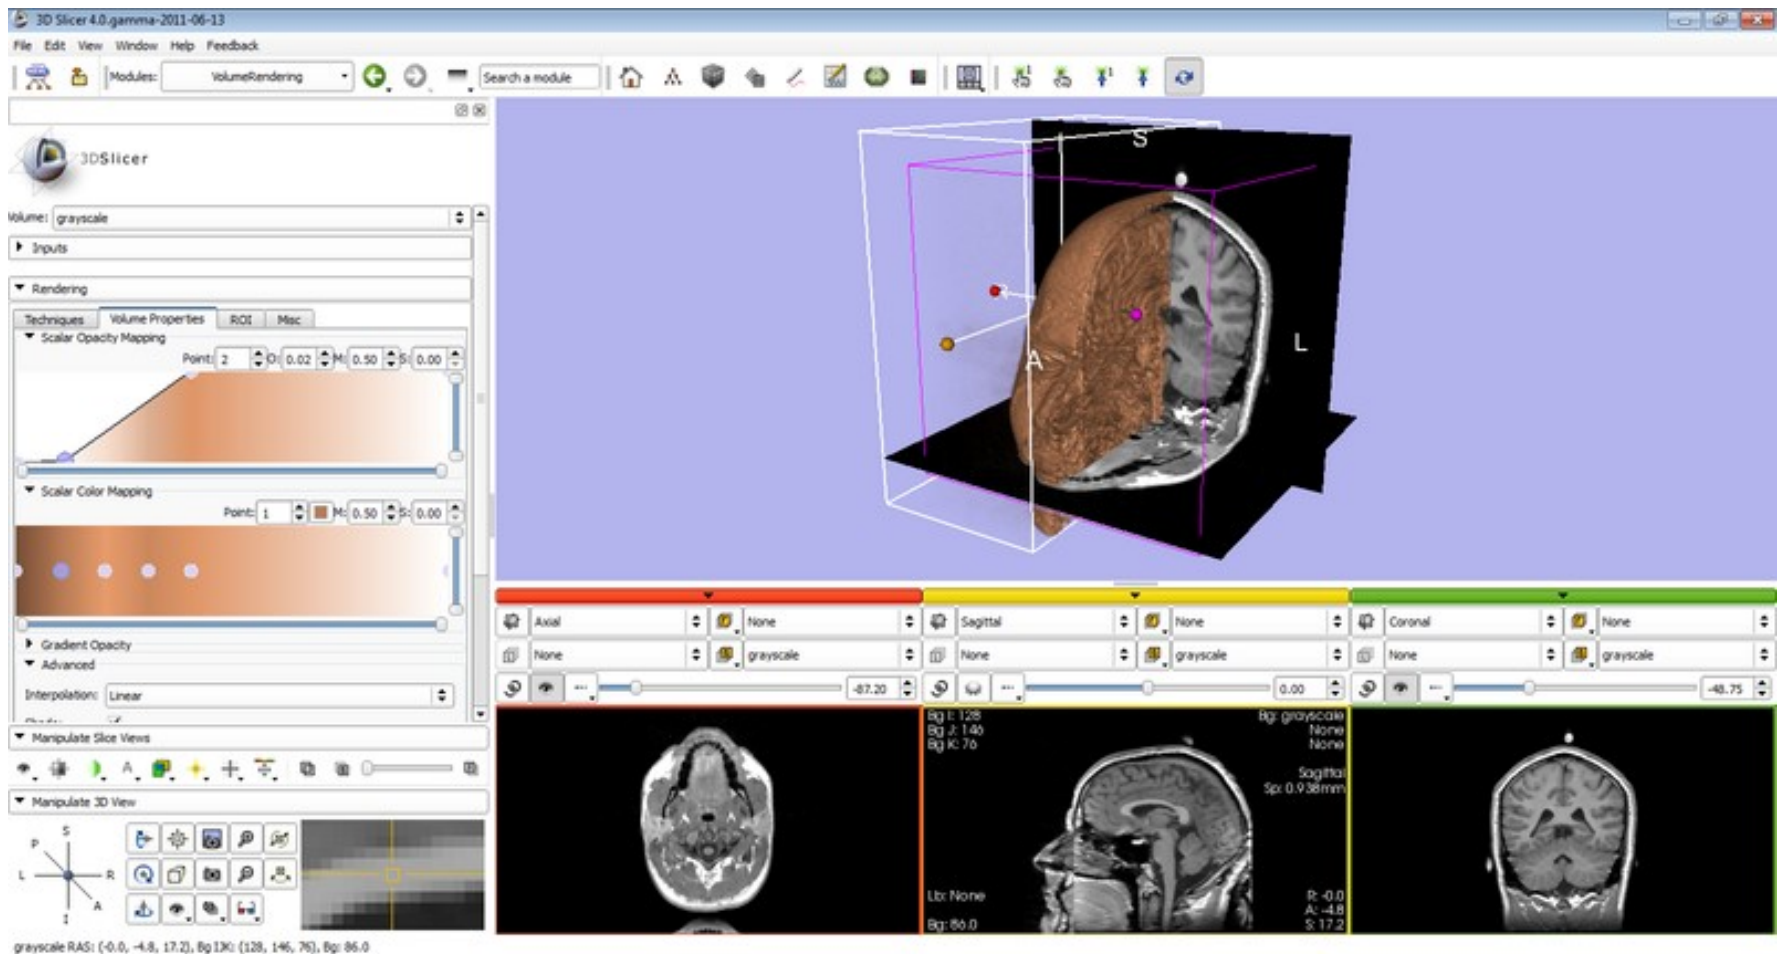


Execution Drivers: Demand-driven

- Launch execution in response to requests for data
- Execution is initiated at the bottom of the pipeline in a sink
 - Request is passed to upstream modules to source
 - Once request reaches a source, it produces data and returns execution back to downstream
 - Called *pull model*
- ex) render request to update GUI



Example: Interactive GUI



Caching Intermediate Values

- Store intermediate results
- Reduce re-run the entire pipeline
 - For example, if a user change the parameter for a specific module using GUI, modules upstream of that module do not need to be executed if the result is cached
- Speed – memory tradeoff



Control

- Centralized
 - Single unit managing the execution of all modules in the pipeline
 - Control unit has links to all modules
 - Complicated to implement, poor scalability, better cache / load balancing
- Distributed
 - Separate unit for each module
 - Use messages to propagate execution
 - Simple to implement, better scalability



Out-of-core Streaming

- Out-of-core
 - Data is too large to fit within internal memory
- Streaming
 - Way to perform out-of-core in pipeline
 - Read data in pieces and let it flow through pipeline




Out-of-core Streaming

- Types of algorithm for streaming
 - Separable (break into pieces)
 - Result invariant (processing order independent)
 - Mappable (should be able to determine which portion of input is required for output)
- Ghost (halo) cells
 - Extra layer of cells at boundary
 - Required for streaming



Example

- Gaussian smoothing
 - Separable, result invariant, mappable?



Example

- Non-local Mean
 - Separable, result invariant, mappable?



$$NL(v)(i) = \sum_{j \in I} w(i, j) v(j)$$


Metadata

- Modern visualization pipelines employ *metadata* – a brief description of actual data
- Metadata can flow through the pipeline independent from the data
- Regions, time, contracts,...




Regions


- How to split, where to modify
- Extents (index range), pieces (collection of cells), blocks (logical domain decomposition)



(a) Update Information



(b) Update Region



(c) Update Data

Three pipeline passes for using regional data




Time

- Executing pipeline over a sequence of time steps
- Add time dimension to the region metadata
- Source declares all available time steps
- Each filter can augment time during the update information pass



Contracts

- General data structure for Impact of a filter
 - Regions, variables, time step, operating restrictions (support of streaming or requiring of ghost cell)
 - Filters declare their impact by modifying a contract object



v0, v1... : contract versions

LB: load balancer



Example: Prioritized streaming

- Change the order of processing based on priority
- Priority metric
 - Close proximity to the viewer in 3D
 - Least likely to be culled
 - Regions with scalar values in an interesting range
 - Regions with more variability



Example: Query-driven visualization

- Analyze a large data set by identifying “interesting” data that matches some specified criteria
- Need to quickly load small selections of data with arbitrary specification
- Required technology
 - File indexing
 - The pipeline source must be able to identify where the pertinent data is located without reading the entire file (e.g., tree-based)
 - Query language
 - E.g., Boolean expression (all regions where (temperature > 1000K) AND (70kPa < pressure < 90kPa))
 - Pipeline metadata mechanism



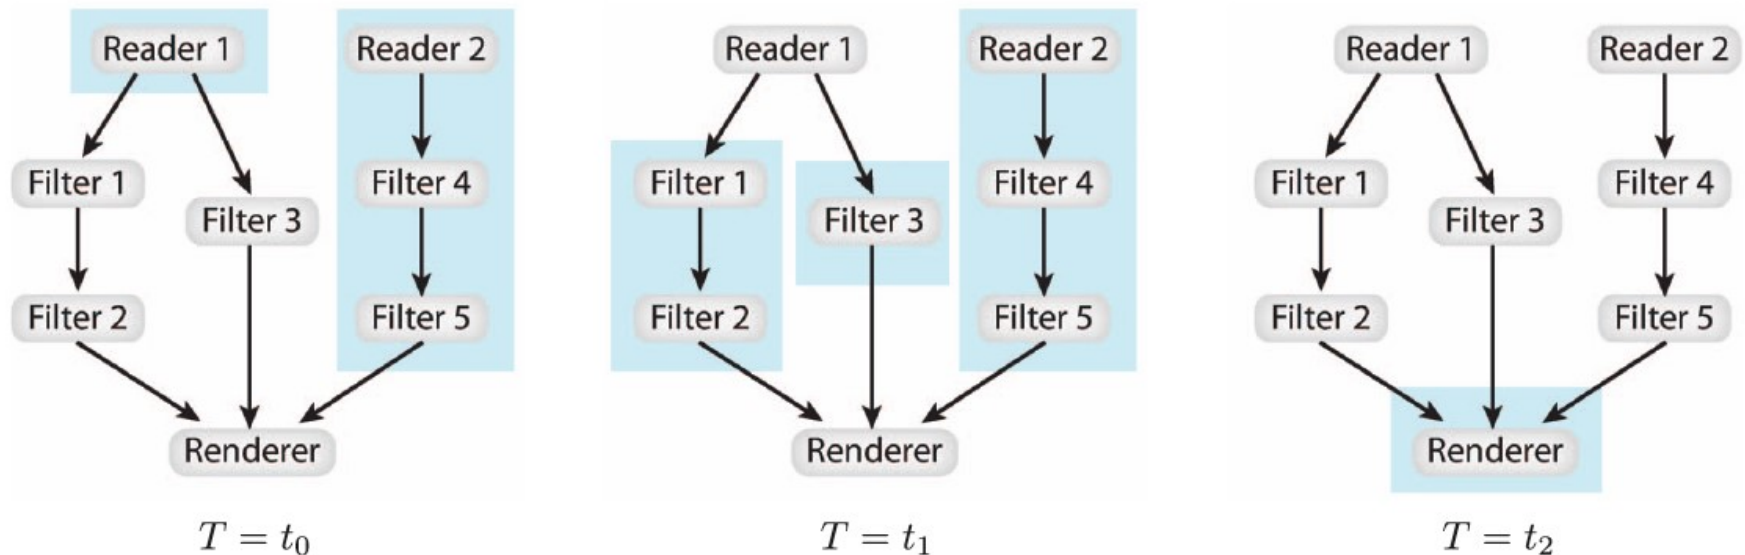
Parallel Execution

- Execute different modules in the pipeline concurrently
- Types of parallel execution
 - Task parallelism
 - Pipeline parallelism
 - Data parallelism



Task Parallelism


- Identify independent portions of the pipeline
 - Use task dependency
- Load balancing is an issue




Mapping Technique for Load Balancing

- Tasks must be mapped to processes
 - Minimizing overhead is the goal
- Overhead
 - Communication
 - Idling

9-12 can start after 1-8 are done



(a)



(b)

Same load balance but (b) has more idling



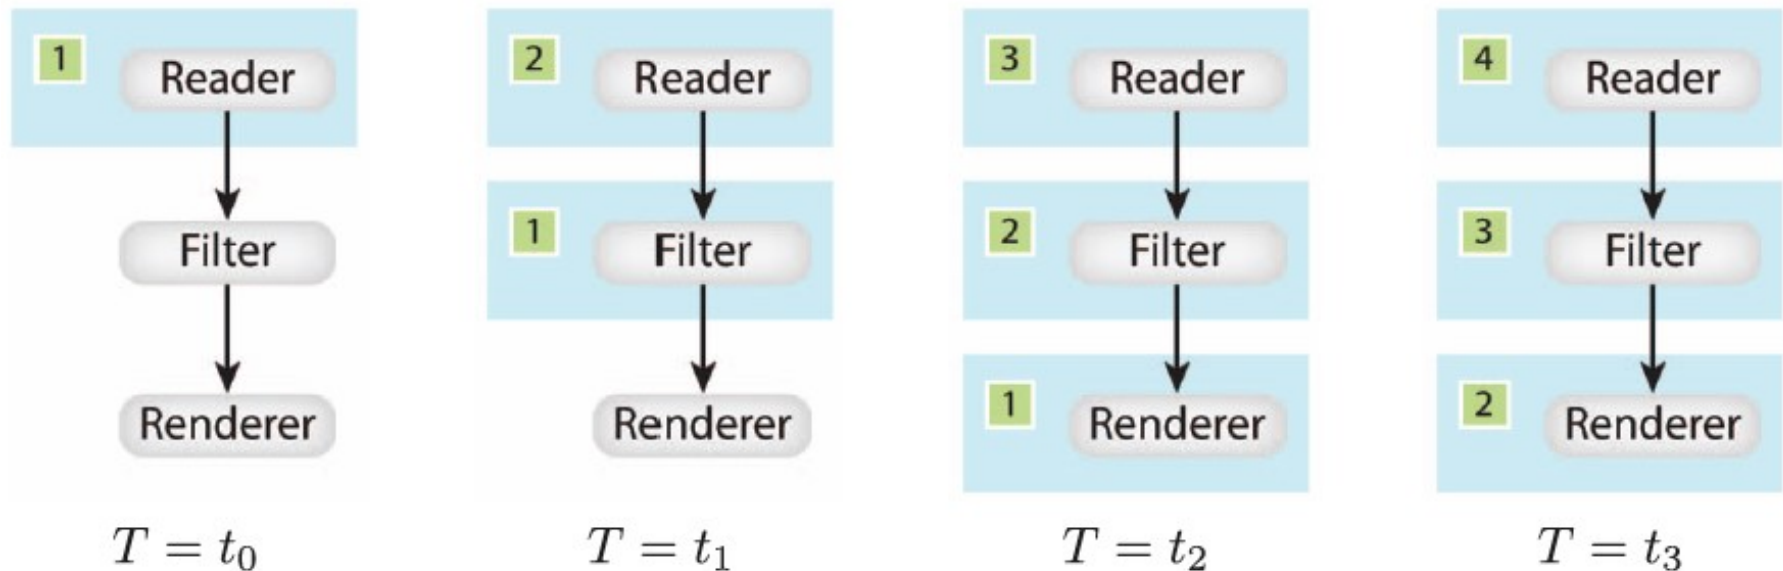
Mapping Techniques

- Static Mapping
 - Tasks are mapped to processes a-priori
 - Need a good estimate of the size of each task
 - The mapping problem can be computationally difficult for non-uniform tasks
- Dynamic Mapping
 - Tasks are mapped to processes at runtime
 - Tasks are generated at runtime, or that their sizes are not known



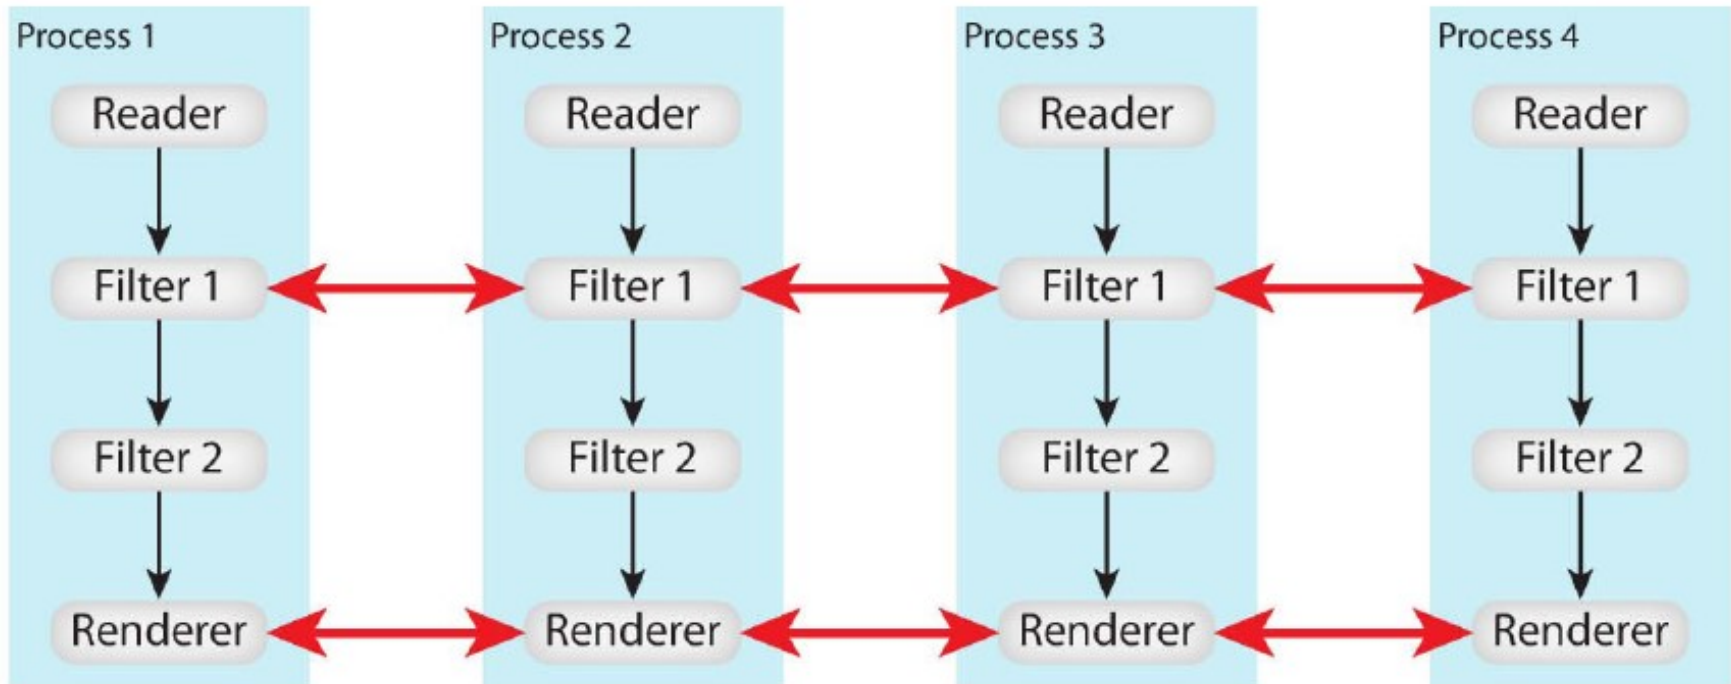
Pipeline Parallelism

- Streaming data in pieces, execute different modules of the pipeline concurrently




Data Parallelism

- Partition the input data, replicate pipeline for each piece



Sort-last Rendering

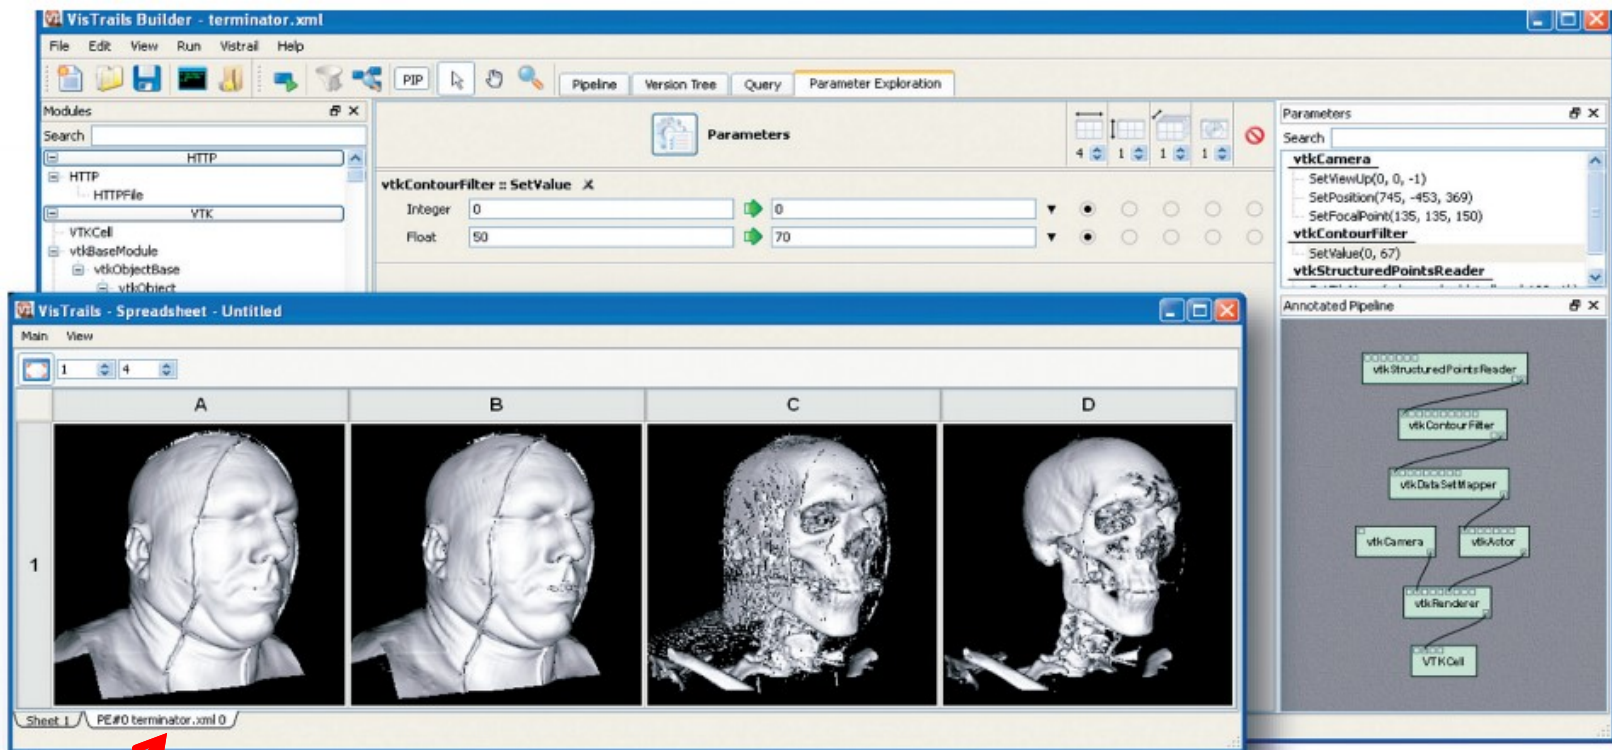
- Fits well to data parallelism
- Each piece is rendered independently
- Merge results at the end



Sort: depth test

Provenance

- Model exploration as transformations to the visualization pipeline



Spreadsheet

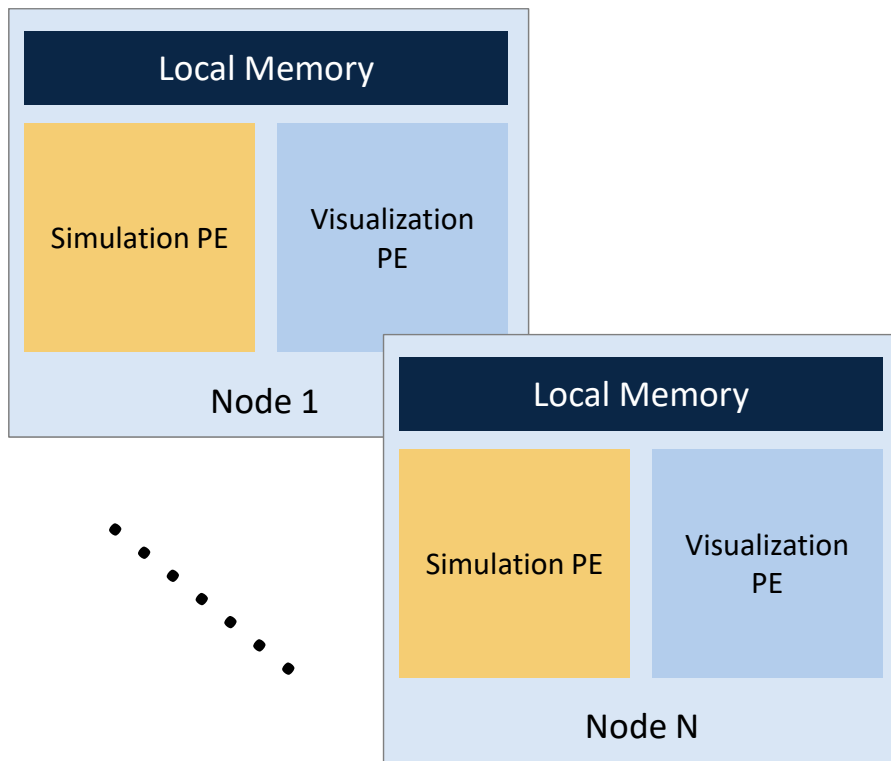
Parameter exploration using VisTrail



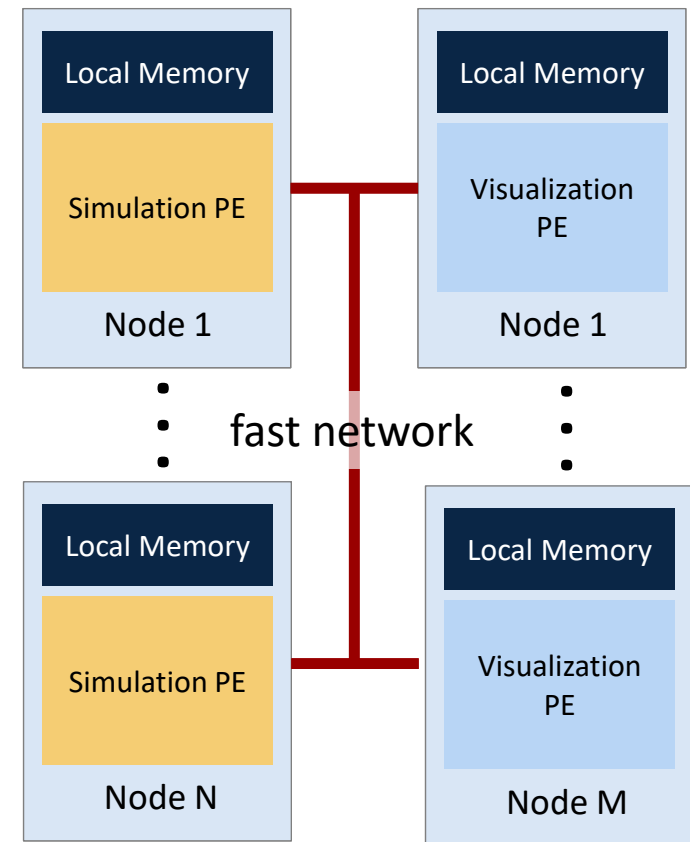
KOREA
UNIVERSITY

In Situ Visualization

- Simulation and visualization run concurrently



Tightly Coupled N nodes

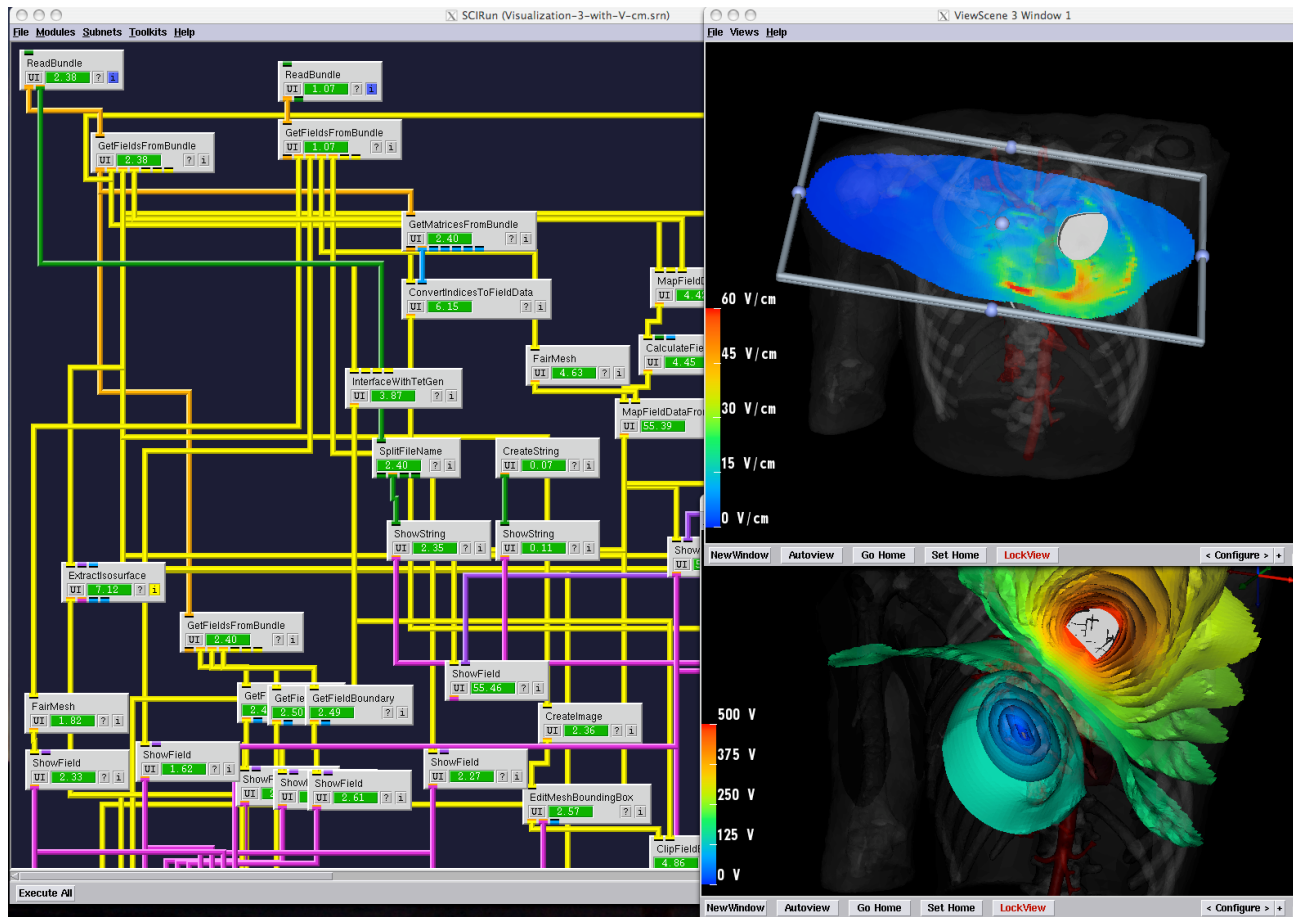


N simulation nodes,
M visualization nodes



Dataflow Visualization Systems

- SCIRun



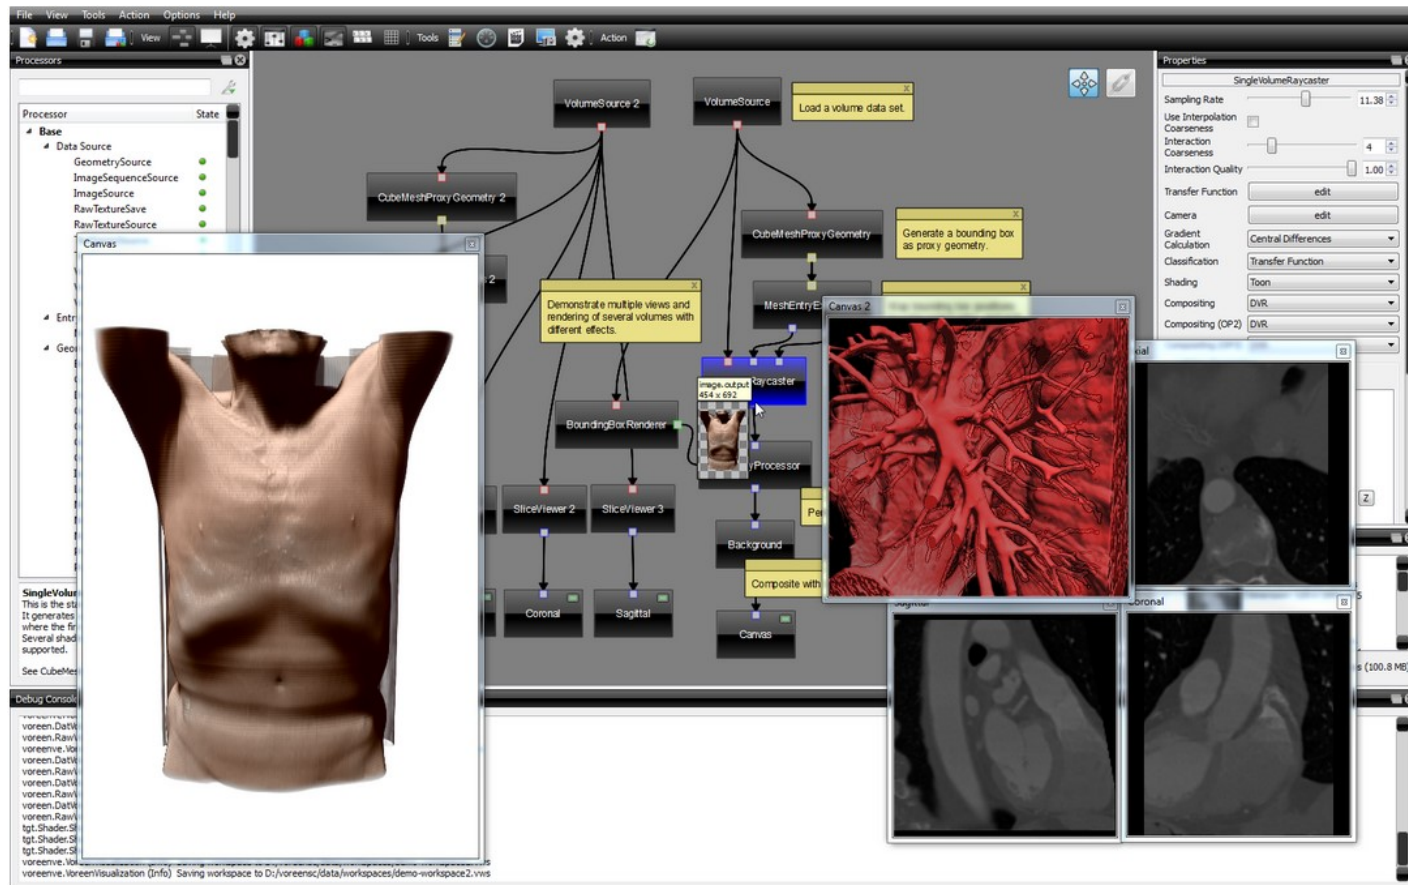
Modeling, simulation, visualization in one package



KOREA
UNIVERSITY

Dataflow Visualization Systems

- Voreen


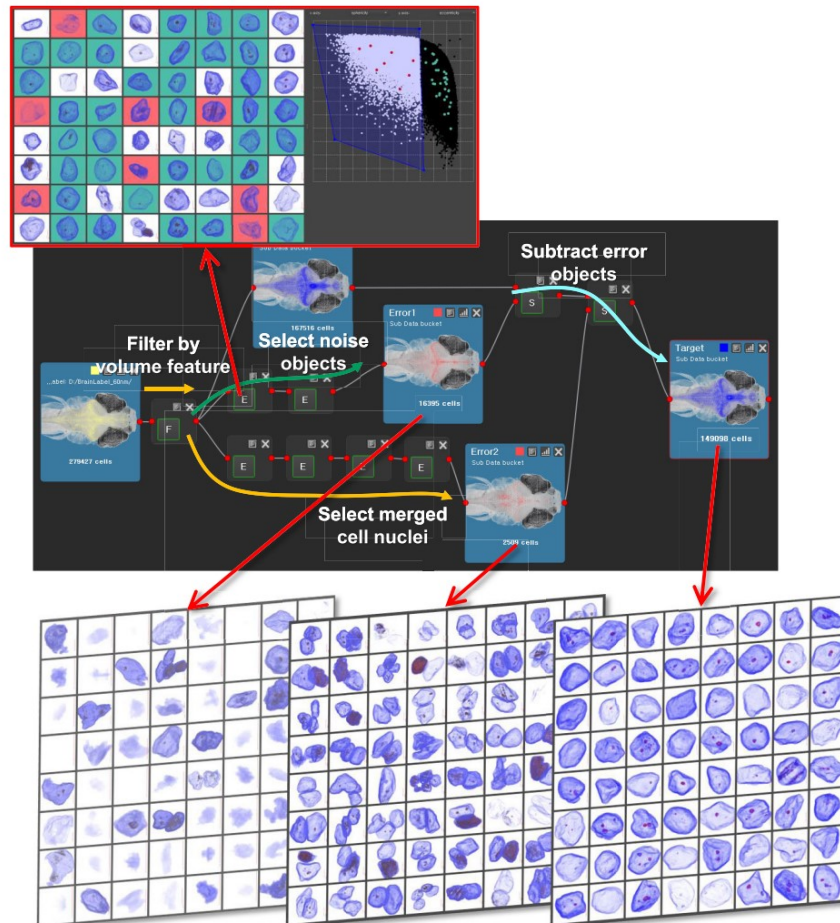


Volume Rendering Engine



Dataflow Visualization Systems

- ZeVis: brain cell morphology analysis



Visualization Pipeline Alternative

- Function field model
 - $f(x) = y$, x : location, y : value
 - Continuous field, derived fields
- MapReduce
 - Simple programming & execution model for distributed processing
- Fine-grained data parallelism
 - Many concurrent threads, multi-CPU or GPU
- Domain-specific languages



Questions?



KOREA
UNIVERSITY